

M5L1 Solutions

May 1, 2023

Exercise 7

Perform cluster analysis on DJIA Index components using K-Means and build an equal weight portfolio from the selected stocks. Retrieve the list of DJIA components; identify and clean any missing data points. Cluster stocks based on weekly ATR and compare it with the original dataset.

Solutions

K-Means Clustering

Clustering is a branch of unsupervised machine learning models that seeks to learn from the properties of the data by identifying groups or clusters in the dataset.

The k -means algorithm searches for a predetermined number of clusters within an unlabeled dataset and is based on the assumptions that the optimal cluster will have cluster center and each point is closer to its own cluster center than to other cluster centers.

```
[1]: # Ignore warnings
import warnings
warnings.filterwarnings('ignore')

# Import Libraries
import pandas as pd
import numpy as np
import pyfolio as pf

from kneed import KneeLocator
import matplotlib.pyplot as plt

from sklearn.cluster import KMeans
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import MinMaxScaler
```

```
[2]: # Load the pre-saved data dict
ohlcv = np.load('data/dow_ohlcv.npy', allow_pickle='TRUE').item()
ohlcv['MMM'].head()
```

```
[2]:
```

	Open	High	Low	Close	Volume
Date					
2009-12-31	62.906018	63.123738	61.967570	62.065166	2049800

2010-01-04	62.380479	62.650753	62.065162	62.327927	3043700
2010-01-05	62.162779	62.485606	61.336941	61.937550	2847000
2010-01-06	62.973588	63.514135	62.695811	62.815929	5268500
2010-01-07	62.553137	62.883472	61.652230	62.860950	4470100

```
[3]: # List of DJIA stocks
dow_stocks = ['MMM', 'AXP', 'AMGN', 'AAPL', 'BA', 'CAT', 'CVX', 'CSCO', 'KO',
↪ 'DOW', 'GS', 'HD', 'HON', 'IBM', 'INTC', 'JNJ', 'JPM', 'MCD', 'MRK', 'MSFT',
↪ 'NKE', 'PG', 'CRM', 'TRV', 'UNH', 'VZ', 'V', 'WBA', 'WMT', 'DIS']
```

```
[4]: # Function to calculate average true range
def ATR(df,n):
    "function to calculate Average True Range"
    df = df.copy()

    df['H-L'] = abs(df['High']-df['Low'])
    df['H-PC'] = abs(df['High']-df['Close'].shift(1))
    df['L-PC'] = abs(df['Low']-df['Close'].shift(1))

    df['TR'] = df[['H-L', 'H-PC', 'L-PC']].max(axis=1,skipna=False)
    df['ATR'] = df['TR'].rolling(n).mean()

    df2 = df.drop(['H-L', 'H-PC', 'L-PC'],axis=1)
    return df2['ATR']
```

```
[5]: # Add ATR for each stocks
for symbol in dow_stocks:
    ohlc[symbol]['ATR'] = ATR(ohlc[symbol],21)
```

```
[6]: # Subsume into dataframe
df = pd.DataFrame({symbol: ohlc[symbol]['ATR'] for symbol in dow_stocks})

# Check for missing values
df.isnull().sum()
```

```
[6]: MMM      22
     AXP      22
     AMGN     22
     AAPL     22
     BA       22
     CAT      43
     CVX      22
     CSCO     22
     KO       22
     DOW     2340
     GS       22
     HD       22
```

```

HON      22
IBM      22
INTC     22
JNJ      22
JPM      22
MCD      22
MRK      22
MSFT     22
NKE      22
PG       22
CRM      22
TRV      22
UNH      22
VZ       22
V        22
WBA      22
WMT      22
DIS      22
dtype: int64

```

```

[7]: # Fill forward the missing values and drop DOW company from the list
df.fillna(method='bfill', axis=0, inplace=True)
df.drop(['DOW'], axis=1, inplace=True)

```

```

[8]: # Resample to a weekly timeframe for cluster analysis
px = df.resample('W-FRI').mean()
px = px.T

# Check output
px.head(2)

```

```

[8]: Date    2010-01-01    2010-01-08    2010-01-15    2010-01-22    2010-01-29    2010-02-05    \
MMM         1.113625     1.113625     1.113625     1.113625     1.113625     1.153880
AXP         1.035185     1.035185     1.035185     1.035185     1.035185     1.046649

Date    2010-02-12    2010-02-19    2010-02-26    2010-03-05    ...    2020-07-31    \
MMM         1.233820     1.308653     1.232967     1.071002    ...         3.304308
AXP         1.079262     1.101860     0.908801     0.774188    ...         2.843043

Date    2020-08-07    2020-08-14    2020-08-21    2020-08-28    2020-09-04    2020-09-11    \
MMM         3.337810     3.339979     3.176965     2.863492     2.869821     3.166674
AXP         2.639048     2.661429     2.715334     2.638858     2.682953     2.666072

Date    2020-09-18    2020-09-25    2020-10-02
MMM         3.413893     4.169714     4.322380
AXP         2.634953     3.009524     3.067619

```

[2 rows x 562 columns]

Elbow Plot The number of clusters is a hyperparameter to clustering models and choose the optimal number of clusters is critical for the model. We identify the elbow point programmatically for this exercise.


```
[9]: scaler = MinMaxScaler()
scaled_px = scaler.fit_transform(px)
```

```
[10]: # Get the inertia
sse = []
for k in range(1, 30):
    kmeans = KMeans(n_clusters=k)
    kmeans.fit(scaled_px)
    sse.append(kmeans.inertia_)

# Knee Locator
kl = KneeLocator(range(1, 30), sse, curve="convex", direction="decreasing")
kl.elbow
```

```
[10]: 6
```

```
[11]: # Raw data and knee.
kl.plot_knee()
```



Build Clusters

We will now fit the cluster model.

```
[12]: # Build clusters
model = KMeans(n_clusters=6)
model.fit(scaled_px)

labels = model.predict(scaled_px)
labels
```

```
[12]: array([1, 0, 3, 4, 5, 1, 1, 4, 4, 2, 0, 0, 1, 4, 0, 0, 0, 4, 0, 4, 0, 1,
        0, 5, 4, 0, 0, 0, 0], dtype=int32)
```

```
[13]: # Remove DOW from the list
companies = dow_stocks
```

```
companies.remove('DOW')
```

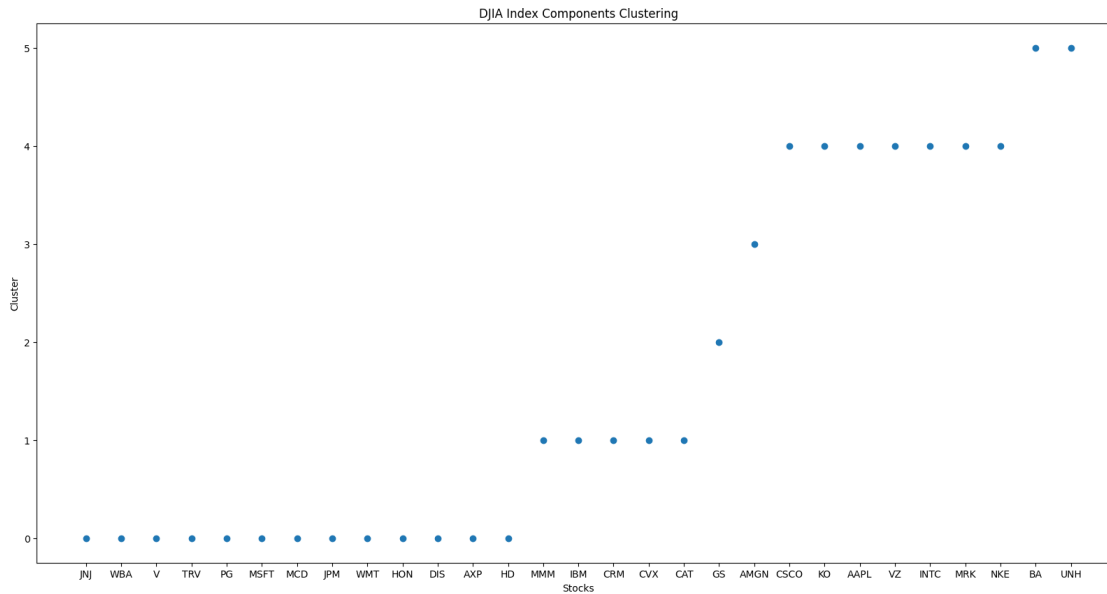
```
[14]: df1 = pd.DataFrame({'Cluster': labels,
                        'Companies': companies,
                        'ATR': px.mean(axis=1),
                        }).sort_values(by=['Cluster'], axis = 0)

df1 = df1.reset_index(drop=True)
df1
```

```
[14]:
```

	Cluster	Companies	ATR
0	0	JNJ	1.306309
1	0	WBA	1.119627
2	0	V	1.501398
3	0	TRV	1.423909
4	0	PG	1.030430
5	0	MSFT	1.290491
6	0	MCD	1.650644
7	0	JPM	1.310115
8	0	WMT	1.116245
9	0	HON	1.606682
10	0	DIS	1.486680
11	0	AXP	1.424071
12	0	HD	1.987495
13	1	MMM	2.101291
14	1	IBM	2.095919
15	1	CRM	2.199421
16	1	CVX	1.630457
17	1	CAT	2.079966
18	2	GS	3.640206
19	3	AMGN	2.664776
20	4	CSCO	0.538751
21	4	KO	0.493262
22	4	AAPL	0.680922
23	4	VZ	0.606684
24	4	INTC	0.695828
25	4	MRK	0.865679
26	4	NKE	0.969232
27	5	BA	3.791872
28	5	UNH	2.748506

```
[15]: # Plot Clusters
plt.figure(figsize=(20,10))
plt.scatter(df1.Companies, df1.Cluster)
plt.xlabel('Stocks')
plt.ylabel('Cluster')
plt.title('DJIA Index Components Clustering');
```



Portfolio Construction Shortlisting the six stocks from the above clusters (one for each), we will now build a portfolio and compare the returns with all stock portfolio.

```
[16]: # Cluster portfolio stocks
portfolio_stocks = ['GS', 'JNJ', 'UNH', 'CSCO', 'CRM', 'AMGN']
port = pd.DataFrame({symbol: ohlc[symbol]['Close'] for symbol in
    ↳ portfolio_stocks})
port.dropna(inplace=True)

port
```

```
[16]:
```

	GS	JNJ	UNH	CSCO	CRM \
Date					
2009-12-31	144.296997	46.714996	25.778193	18.048286	18.442499
2010-01-04	147.920776	46.910812	26.666224	18.613705	18.705000
2010-01-05	150.535919	46.366856	26.623932	18.530785	18.625000
2010-01-06	148.929138	46.743999	26.886124	18.410154	18.592501
2010-01-07	151.843475	46.410374	27.917923	18.493084	18.510000
...
2020-09-23	186.119995	144.440002	292.140015	37.930145	235.990005
2020-09-24	195.110001	144.669998	292.660004	37.504074	237.550003
2020-09-25	194.949997	145.660004	302.500000	38.098591	242.740005
2020-09-28	199.070007	147.110001	303.230011	38.772377	246.669998
2020-09-29	196.789993	147.059998	304.149994	38.703018	247.449997

	AMGN
Date	

```

2009-12-31    45.404182
2010-01-04    46.327202
2010-01-05    45.925884
2010-01-06    45.580757
2010-01-07    45.163395
...
2020-09-23    242.589996
2020-09-24    240.320007
2020-09-25    243.820007
2020-09-28    247.029999
2020-09-29    248.300003

```

[2705 rows x 6 columns]

```

[17]: # Calculate portfolio returns
portfolio_returns = port.pct_change().fillna(0)
port['Returns'] = portfolio_returns.mean(axis=1)
port.head(2)

```

```

[17]:
      Date      GS      JNJ      UNH      CSC0      CRM      AMGN  \
2009-12-31  144.296997  46.714996  25.778193  18.048286  18.442499  45.404182
2010-01-04  147.920776  46.910812  26.666224  18.613705  18.705000  46.327202

      Returns
Date
2009-12-31  0.000000
2010-01-04  0.021607




```

```

[18]: # Create Tear sheet using pyfolio
pf.create_simple_tear_sheet(port['Returns'])

```

<IPython.core.display.HTML object>



```
[19]: # All stocks portfolio
all_stocks = pd.DataFrame({symbol: ohlc[symbol]['Close'] for symbol in
    ↪companies})
all_stocks.dropna(inplace=True)
all_stocks.head(2)
```

```
[19]:
```

	MMM	AXP	AMGN	AAPL	BA	CAT	\
Date							
2009-12-31	62.065166	34.434761	45.404182	6.503574	41.856789	41.498463	
2010-01-04	62.327927	34.774685	46.327202	6.604801	43.441975	42.634403	

	CVX	CSCO	KO	GS	...	NKE	\
Date					...		
2009-12-31	50.924435	18.048286	18.951757	144.296997	...	12.066024	
2010-01-04	52.293617	18.613705	18.965061	147.920776	...	11.934528	

	PG	CRM	TRV	UNH	VZ	V	\
Date							
2009-12-31	43.431492	18.442499	38.205708	25.778193	18.633041	17.935587	
2010-01-04	43.782478	18.705000	38.167397	26.666224	18.717409	18.075037	

	WBA	WMT	DIS
Date			
2009-12-31	28.350834	40.954620	28.090706
2010-01-04	28.798639	41.552284	27.933924

[2 rows x 29 columns]

```
[20]: # Calculate all stocks portfolio returns
all_stocks_returns = all_stocks.pct_change().fillna(0)
all_stocks['Returns'] = all_stocks_returns.mean(axis=1)
all_stocks.head(2)
```

```
[20]:
```

	MMM	AXP	AMGN	AAPL	BA	CAT	\
Date							
2009-12-31	62.065166	34.434761	45.404182	6.503574	41.856789	41.498463	
2010-01-04	62.327927	34.774685	46.327202	6.604801	43.441975	42.634403	

	CVX	CSCO	KO	GS	...	PG	\
Date					...		
2009-12-31	50.924435	18.048286	18.951757	144.296997	...	43.431492	
2010-01-04	52.293617	18.613705	18.965061	147.920776	...	43.782478	




	CRM	TRV	UNH	VZ	V	WBA	\
Date							
2009-12-31	18.442499	38.205708	25.778193	18.633041	17.935587	28.350834	
2010-01-04	18.705000	38.167397	26.666224	18.717409	18.075037	28.798639	

	WMT	DIS	Returns
Date			
2009-12-31	40.954620	28.090706	0.000000
2010-01-04	41.552284	27.933924	0.013946

[2 rows x 30 columns]

```
[21]: # Create Tear sheet using pyfolio
pf.create_simple_tear_sheet(all_stocks['Returns'])
plt.show()
```

<IPython.core.display.HTML object>



Conclusion The cluster stocks generated a CAGR of 17% with a maximum drawdown of ~30% while the all stocks portfolio consisting of 29 DJIA index stocks generated an annualized return of ~16% with a maximum drawdown of 33%. This study highlight that with 20% of all stocks, we can construct a portfolio that can outperform the all stocks portfolio with an alpha of 59% and an improved sortino ratio.

Note:

1. The data is not treated for in/out sample as the objective here is to showcase the application of clustering methods. Accordingly, the actual results may vary.
2. Arbitrary selection of Cluster stocks can be avoided by adopting a minimum distance measure in stock selection.

References

- [Scikit-learn K-Means Clustering](#)
- [Pyfolio-reloaded](#)
- [Python resources](#)

* * *